

Classification

Classification vs. regression


In supervised learning we try to find a function f , which systematically produces the output values y_m associated with the input values $X_{m,:}$:

$$f(X_{m,:}) \rightarrow y_m$$

Classification

Target variable y : categorical


$$y_m \in \{C_1, C_2, \dots, C_K\}$$





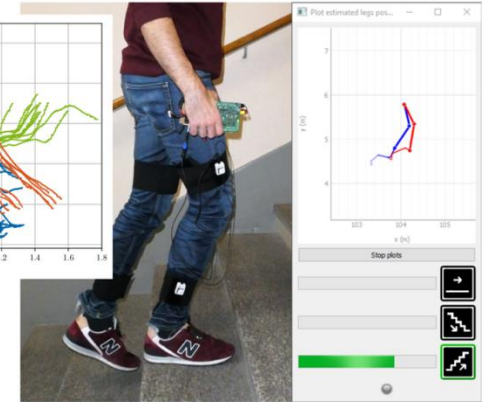


Regression

Target variable y : numerical - continuous

$$y_m \in \mathbb{R}$$



Classification examples



Text classification

Categorize text documents into predefined categories . For example, categorize news into 'sports', 'politics', 'science', etc.

Soft tissue found in T-rex fossil

Find may reveal details about cells and blood vessels

Health may be concern when giving kids cell phones

Wall Street gears up for jobs

Probe finds atmosphere on Saturn moon


Thursday, March 17, 2005 Posted: 11:17 AM EST

LOS ANGELES, California (Reuters) -- The space probe

**Cassini discovered a significant atmosphere around Saturn's
moon Enceladus during two recent passes close by, the Jet**

Character recognition

Identify handwritten characters:
classify each image of a
character into one of 10
categories '0', '1', '2' ...



Types of classification problems

Binary: 2 possible classes

- Structured data

Credit amount	cash	age	purpose	Risk
60'000	25'000	55	«Auto»	
...				

Risk
0
1

- Images



Hot Dog



No Hot Dog

- Text

- Sentiment analysis:

positive

negative

«Great pasta and perfect location!»

«Bad service – never again!»

- Spam detection: spam/ham

Multinomial (multi-class): >2 possible classes



Iris Versicolor

Iris Setosa

Iris Virginica



French

English

German

Multilabel:



Love story

Action

Comedy

Science fiction

Thriller


Classification algorithms

Generative classifiers


e.g. Naive Bayes

$$p(C_c|x) \propto p(C_c) \prod_{n=1}^N p(x_n|C_c)$$


Decision Trees




Support Vector Machines (SVM)



Logistic regression



Neural networks




Logistic Regression

Logistic regression solves a binary classification task

**Logistic regression,
although termed
'regression' it is used
for classification!**

- Given are M training samples $(\mathbf{X}_{m,:}, y_m)$
- Each $\mathbf{X}_{m,:}$ is a N -dimensional **feature vector**
the features (independent variables) can be continuous or discrete
- $y_m \in \{0, 1\}$ are the **discrete labels**.
Binary classification: 0/1, i.e. true/false, positive/negative ...

Why not regression?



The hypothesis in logistic regression

Logistic Function (or "Sigmoid Function"):

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


with $z = \theta_0 + \theta_1 x = \boldsymbol{\theta}^T \mathbf{X}_{m,:}$:

Properties of the sigmoid σ :

- smooth distribution between 0 and 1
- $\sigma(z) = 0.5$
- convenient derivatives:


$$\sigma'(z) = \sigma(z) \cdot (1 - \sigma(z))$$

Hypothesis: $h(\mathbf{X}_{m,:}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}^T \mathbf{X}_{m,:})$




Effect of the parameters in logistic Function

$\sigma(z)$




x

$$\frac{1}{1 + e^{-(0+1 \cdot x)}}$$



x

$$\frac{1}{1 + e^{-(0+20 \cdot x)}}$$



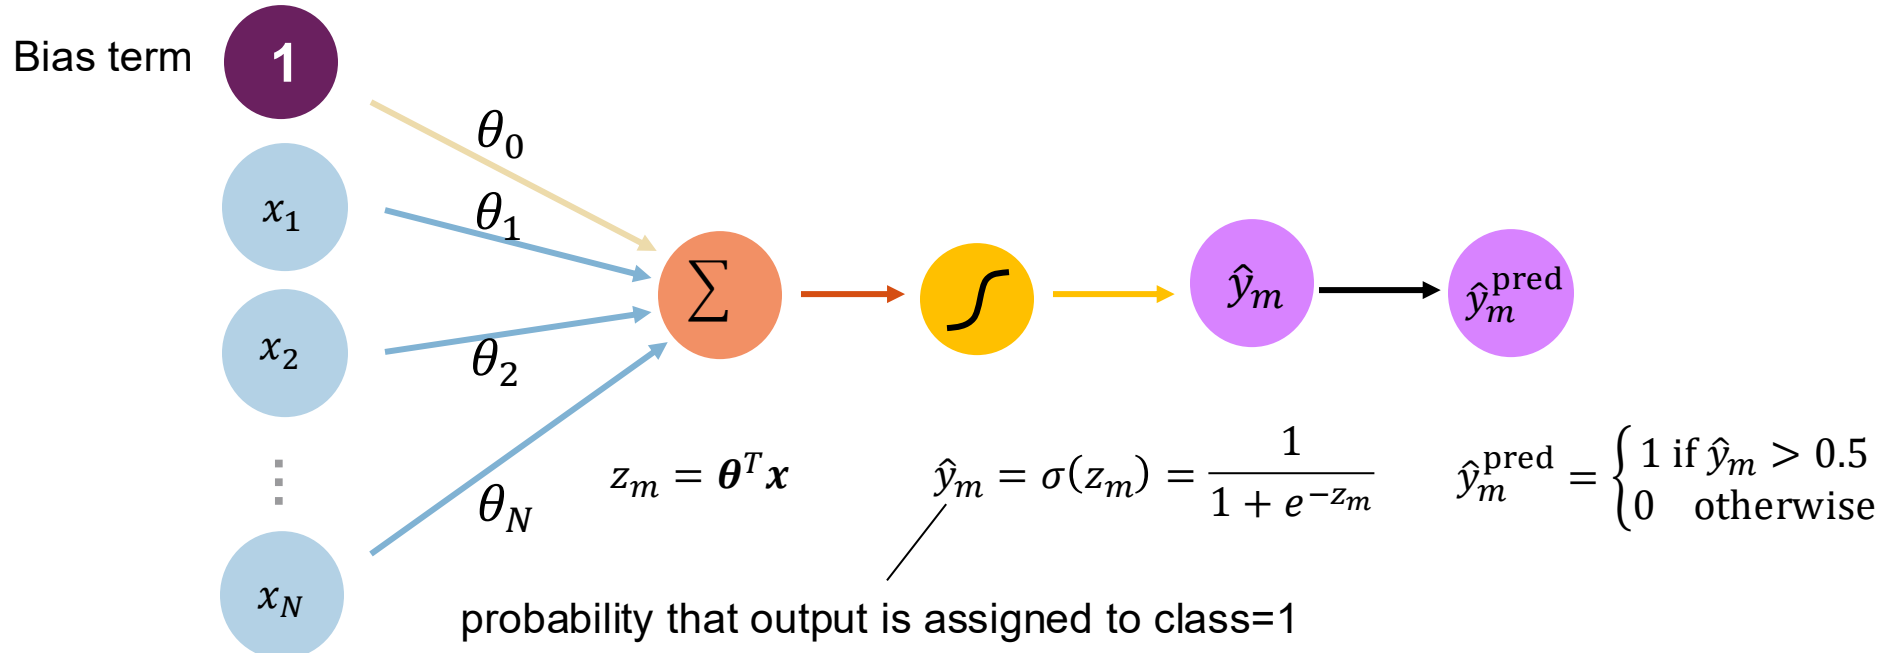
x

$$\frac{1}{1 + e^{-(50+20 \cdot x)}}$$

Logistic regression for one sample, N features

Forward pass on sample $X_{m,:}$ (m -th row in X): x

Input x weights linear combination non-linearity output decision



The loss function in logistic regression

For the N -dimensional training example X_m , the model yields a value $0 < \hat{y}_m < 1$

If the corresponding label y_m is 1, the likelihood of y_m being the positive class according to the model, is given by \hat{y}_m . If y_m is 0, the likelihood of y_m being the class 0 is given by $1 - \hat{y}_m$

This leads to the likelihood of the M (independent) training samples according to the model:

$$L(\theta) = \prod_{i=1}^N (\hat{y}_m)^{y_m} (1 - \hat{y}_m)^{1-y_m}$$

The log likelihood is more practical:

$$\log L(\theta) = \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m)$$

Log is strictly increasing \rightarrow any θ that maximize the log likelihood also maximise the likelihood, and vice versa.

The base of the logarithm is not important, but needs to be applied consistently

The cross entropy loss


The cross-entropy loss (CE) is the negative log-likelihood

$$\mathcal{L}_{CE}(\theta) = - \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m)$$

with y_m : binary $\{0,1\}$, \hat{y}_m
scalar in the interval $[0,1]$

With natural logarithm:

y_m	\hat{y}_m	\mathcal{L}_{CE}
0	0.0001	0.0001
0	0.2	0.2231
0	0.5	0.6931
0	0.8	1.6094
0	0.9999	9.2103
1	0.9999	0.0001
1	0.8	0.2231
1	0.5	0.6931
1	0.2	1.6094
1	0.0001	9.2103



The cost function in logistic regression

Average cross entropy over all training samples:

$$\begin{aligned} J(\boldsymbol{\theta}) &= -\frac{1}{M} \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m) \\ &= -\frac{1}{M} \sum_{m=1}^M y_m \log \left(\frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{X}_{m,:}}} \right) + (1 - y_m) \log \left(1 - \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{X}_{m,:}}} \right) \end{aligned}$$

Linear regression vs. logistic Regression

Linear Regression

Hypothesis: $h(\mathbf{X}_{m,:}, \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{X}_{m,:}$

Cost Function: Average sum of squared residuals

$$\begin{aligned} J(\boldsymbol{\theta}) &= \frac{1}{2M} \sum_{i=1}^M (y_m - \hat{y}_m)^2 \\ &= \frac{1}{2M} \sum_{m=1}^M (y_m - \boldsymbol{\theta}^T \mathbf{X}_{m,:})^2 \end{aligned}$$

Gradient of the cost function with respect to the parameters:

$$\frac{\partial J(\boldsymbol{\theta})}{\partial \theta_n} = \frac{1}{M} \sum_{m=1}^M (y_m - \hat{y}_m)(-x_{mn})$$

Logistic Regression

$h(\mathbf{X}_{m,:}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}^T \mathbf{X}_{m,:})$

Average cross entropy

$$\begin{aligned} J(\boldsymbol{\theta}) &= -\frac{1}{M} \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m) \\ &= -\frac{1}{M} \sum_{m=1}^M y_m \log \left(\frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{X}_{m,:}}} \right) + (1 - y_m) \log \left(1 - \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{X}_{m,:}}} \right) \end{aligned}$$

$$\frac{\partial J(\boldsymbol{\theta})}{\partial \theta_n} = \frac{1}{M} \sum_{m=1}^M (y_m - \hat{y}_m)(-x_{mn})$$


Gradient Descent for Logistic Regression

Cost function: $J(\boldsymbol{\theta}) = -\frac{1}{M} \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m)$

This leads¹ to the following update rule in Gradient Descent
Repeat until convergence:

$$\theta_n = \theta_n - \alpha \frac{\partial}{\partial \theta_n} J(\boldsymbol{\theta}) \rightarrow \theta_n = \theta_n - \alpha \frac{1}{M} \sum_{m=1}^M (y_m - \hat{y}_m)(-x_{mn})$$

This is similar to the Least Mean Squares Update Rule in Linear Regression, except that here we use a **non-linear function** in the hypothesis to compute the outputs.




Logistic regression is a linear model for binary classification problems


It learns linear decision boundaries between two classes by minimising the cross entropy loss

$$J(\theta) = -\frac{1}{M} \sum_{m=1}^M y_m \log \hat{y}_m + (1 - y_m) \log(1 - \hat{y}_m)$$

with respect to the model parameters θ using gradient descent




It fails with data, that is not linearly separable



Multinomial Classification

Multinomial (Multi-class) Classification


- **Binary vs multinomial classification**



- **Examples of multinomial classification**
 - Visual digit recognition in OCR
 - Obstacle recognition in autonomous driving

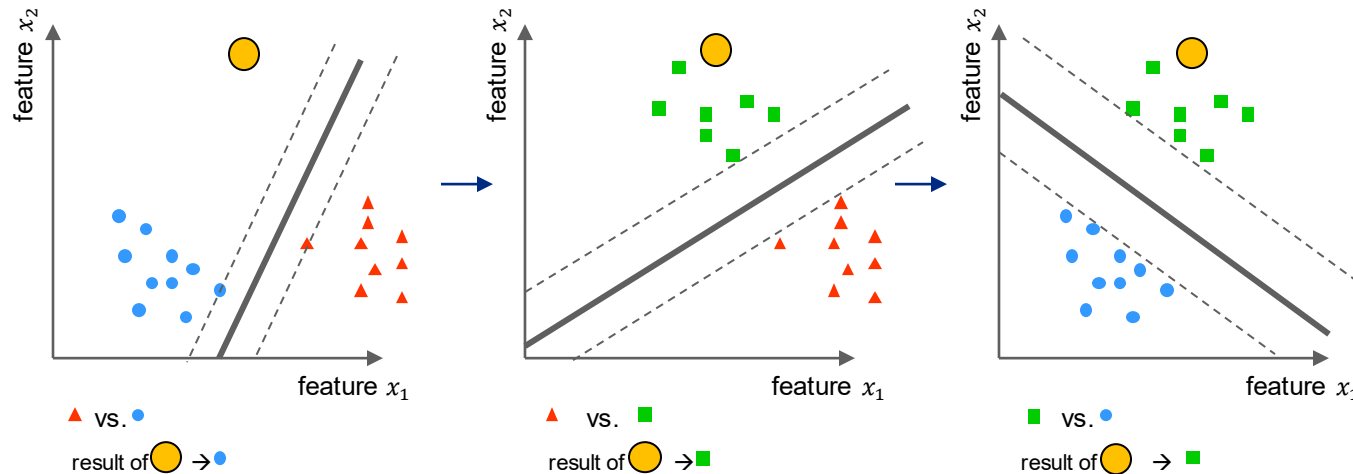
One-vs-Rest approach for multinomial classification

- C classes
- C binary classification problem: each class vs all the others
- C probabilities p_c
- Choose the highest probability to predict the class.



One vs. One

- assume C classes and learn all $C(C - 1)/2$ pairwise comparisons
- classify unknown instance by majority vote among all pairwise comparisons



Softmax regression

The labels are now also encoded in matrix form

Example with $M = 4$ samples and $C = 3$ classes $\{1,2,3\}$:

$$\mathbf{y}_{M \times 1} := \begin{pmatrix} 1 \\ 1 \\ 3 \\ 1 \end{pmatrix} \xrightarrow{\text{one-hot encoding}} \mathbf{Y}_{M \times C} := \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{pmatrix}$$

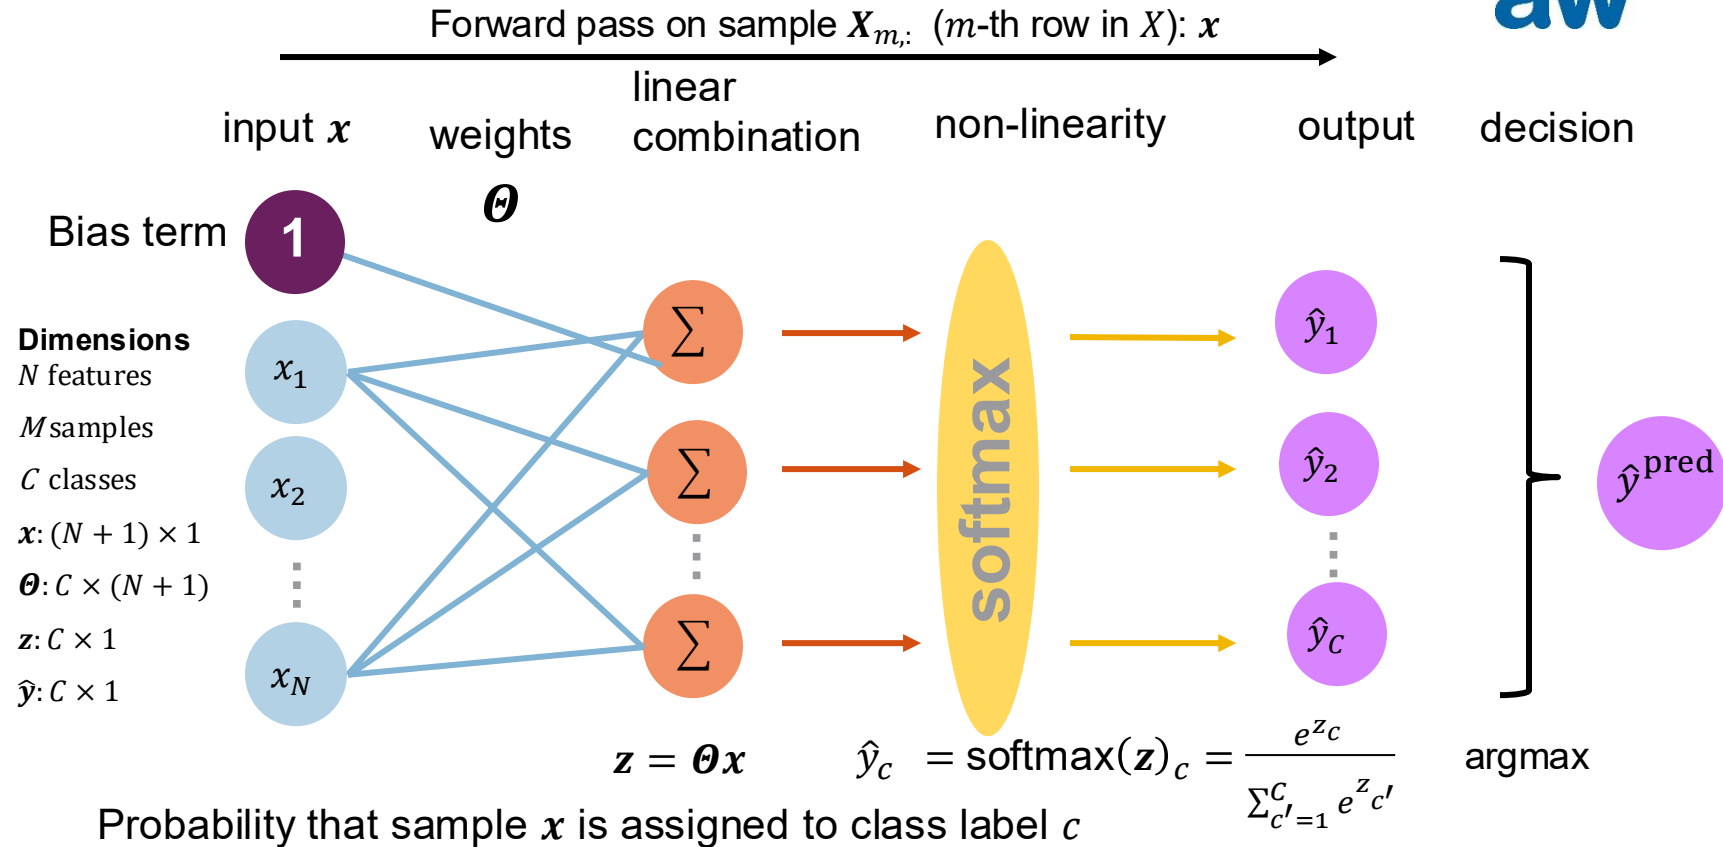
M training samples, C classes

$$\mathbf{Y}_{M \times C} := \begin{pmatrix} y_{11} & y_{12} & \cdots & y_{1C} \\ y_{21} & y_{22} & \cdots & y_{2C} \\ \vdots & \vdots & \ddots & \vdots \\ y_{M1} & y_{M2} & \cdots & y_{MC} \end{pmatrix}$$

m -th row (encoded as a column vector)

$$\mathbf{Y}_{m,:} = \begin{pmatrix} y_{m1} \\ y_{m2} \\ \vdots \\ y_{mC} \end{pmatrix} \text{ is the one-hot encoding of the label for sample } m, \text{ so } y_{mc} = \mathbb{I}(y_m = c)$$

Softmax regression



The learning in softmax regression


The loss function for multinomial logistic regression generalises the loss function for binary logistic regression from 2 to C classes:

$$\begin{aligned}\mathcal{L}_{\text{CE}}(\boldsymbol{\theta}) &= -\frac{1}{M} \sum_{m=1}^M \sum_{c=1}^C y_{mc} \log \hat{y}_{mc} \\ &= -\frac{1}{M} \sum_{m=1}^M y_{mk} \log \hat{y}_{mk} \quad (\text{where } k \text{ is the correct class index for sample } m) \\ &= -\frac{1}{M} \sum_{m=1}^M 1 \cdot \log \frac{\exp(z_k)}{\sum_{c=1}^C \exp(z_c)} = -\frac{1}{M} \sum_{m=1}^M \log \frac{\exp(\boldsymbol{\theta}_{k,:}^T \mathbf{X}_{m,:})}{\sum_{c=1}^C \exp(\boldsymbol{\theta}_{c,:}^T \mathbf{X}_{m,:})}\end{aligned}$$

→ Minimisation by gradient descent


Softmax regression for multinomial classification

Example: 3 Classes sampled from a multivariate normal distribution



Softmax Logistic Regression still learns linear decision boundaries


Softmax regression fails non non-linear problems



Softmax vs One-vs-rest

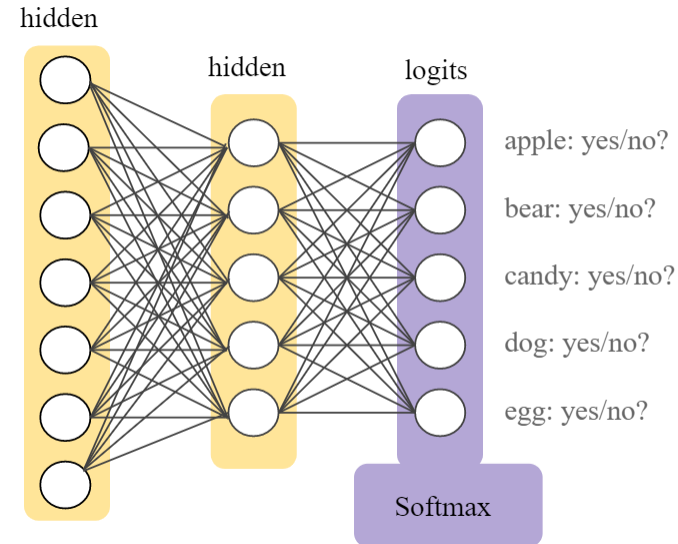
One-vs-rest

Each class has its own model



Softmax


One model for all classes



Nearest Neighbor Classification

Instance Based Classification: K-Nearest Neighbor (KNN)

- The (input,output)-pairs of the training data $\{(X_{m,:}, y_m) | 1 \leq m \leq M\}$ is saved in a search structure
- k : user defined hyperparameter
- To label an **unknown data point**, the **most frequent label among the k closest training samples** is determined → **red**
- Requires a distance metric!
 - E.g. cosine distance





2 Classes: $y_m \in \{\text{red}, \text{blue}\}$

11 training samples: $M = 11$

Evaluation metrics for classification








Can we estimate quantitatively how well the model generalises beyond the training data?

General for supervised machine learning: Split the available data into a training and independent test set



Estimate the generalisation error based on the independent test set

Classification Error and Accuracy - example

inputs x_m	expected outputs y_m	predicted outputs \hat{y}_m	$\mathbb{I}(\hat{y}_m \neq y_m)$
	1	1	
	0	0	
	1	1	
	1	0	
	1	0	
	1	0	
	0	1	

$M = ?$

$$\text{Error} = \frac{1}{M} \sum_{m=1}^M \mathbb{I}(\hat{y}_m \neq y_m)$$

=

$$\text{Accuracy} = 1 - \text{Error} =$$

Classification error and accuracy

The misclassification error for one sample x_m : $\mathbb{I}(\hat{y}_m \neq y_m) = \begin{cases} 1 & \text{if } \hat{y}_m \neq y_m \\ 0 & \text{otherwise} \end{cases}$

delivers a loss of 1 if the predicted value \hat{y}_m does not agree with the expected value y_m

The misclassification rate (standard error) is the average over all samples:

$$\text{Error} = \frac{1}{M} \sum_{m=1}^M \mathbb{I}(\hat{y}_m \neq y_m)$$

$$\text{Accuracy} = 1 - \text{Error}$$

Classification error and accuracy do not distinguish between the different types of errors – some might be more severe than others.

Types of classification errors

Binary classification problem: 0/1

expected \ predicted		
	Not Hot Dog (0)	Hot Dog (1)
Not Hot Dog (0)	<i>TN</i> 	<i>FP</i> 
Hot Dog (1)	<i>FN</i> 	<i>TP</i> 

- True Positives (*TP* - hit): The actual class was 1 (true) and the predicted is also 1 (true)
- True Negatives (*TN*): The actual class was 0 (false) and the predicted is also 0 (false)
- False Positives (*FP* – false alarm, type-I error): The actual class was 0 (false) and the predicted is 1 (true)
- False Negatives (*FN* – miss, type-II error): The actual class was 1 (true) and the predicted is 0 (false)

Types of classification errors

TN




FP



FN




TP




Depending on the use case the types of errors in classification have different impact. This has to be reflected in the selection of the evaluation metrics.

Evaluation metrics for classification


- Accuracy and Error
 - $\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$
 - $\text{Error} = 1 - \text{Accuracy}$
 - Do not take into account different «costs» of errors
- Recall
 - $\text{Recall} = \frac{TP}{TP+FN}$ penalises FN , but ignores FP
 - Fraction of correctly predicted positive elements in relation to all positive elements (relevant)
- Precision
 - $\text{Precision} = \frac{TP}{TP+FP}$ penalises FP , but ignores FN
 - Fraction of correctly predicted positive elements in relation to all positive predicted elements (selected)




How many selected items are relevant?



How many relevant items are selected?



ROC Curve and Precision-Recall Curve Plot





False Positive Rate

$$FP / \#neg = FP / (TN + FP)$$


True Positive Rate (Recall)

$$TP / \#pos = TP / (TP + FN)$$




ROC Curve & PR curve

ROC Curve



PR curve



Evaluation metrics for classification: F-Score

Combination of recall and precision

Definition of the general F_β -Score:

$$F_\beta = (1 + \beta^2) * \frac{\text{Precision} \cdot \text{Recall}}{(\beta^2 \cdot \text{Precision}) + \text{Recall}}$$

F_1 -Score:

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

A smaller β value, such as 0.5, gives more weight to precision and less to recall, whereas a larger β value, such as 2.0, gives less weight to precision and more weight to recall in the calculation of the F -score

True Class	truck	923	9	1	39	2			20	3	3	92.3%	7.7%
	ship	10	931	1	10	2		1	37	4	4	93.1%	6.9%
	horse	4	2	915	1	3	22	17	9	14	13	91.5%	8.5%
	automobile	15	5	1	972				5	2		97.2%	2.8%
	frog		1	1		943	3	4	5	16	27	94.3%	5.7%
	deer	1	2	14	1	14	898	13	5	28	24	89.8%	10.2%
	dog	3		17	2	13	18	801	7	28	111	80.1%	19.9%
	airplane	6	23	5	4	5	4	1	923	21	8	92.3%	7.7%
	bird	3	4	5	2	17	13	8	26	892	30	89.2%	10.8%
	cat	7	5	12	4	30	24	48	12	32	826	82.6%	17.4%

95.0%	94.8%	94.1%	93.9%	91.6%	91.4%	89.7%	88.0%	85.8%	79.0%
5.0%	5.2%	5.9%	6.1%	8.4%	8.6%	10.3%	12.0%	14.2%	21.0%

truck	ship	horse	automobile	frog	deer	dog	airplane	bird	cat
-------	------	-------	------------	------	------	-----	----------	------	-----

Predicted Class

Evaluation metrics for multi-class classification

Medical diagnosis in $C = 3$ classes: {Healthy, Cold, Flue}

100 test results:

		Predicted label			
True label	c	Healthy	Cold	Flue	Support of c
	Healthy	60	8	2	70
	Cold	4	12	4	20
	Flue	0	2	8	10
					100

Average metrics:

$$\text{Metric}_{\text{macro}} = \frac{\sum_{c=1}^C \text{Metric}_c}{C}$$

$$\text{Metric}_{\text{weighted}} = \sum_{c=1}^C \frac{\text{Metric}_c}{\text{Support}_c}$$

Metrics per class:

$$\text{Precision}_c = \frac{TP_c}{TP_c + FP_c}$$

$$\text{Recall}_c = \frac{TP_c}{TP_c + FN_c}$$

$$F1_c = 2 \frac{\text{Precision}_c \cdot \text{Recall}_c}{\text{Precision}_c + \text{Recall}_c}$$

$$\text{Accuracy} = \frac{\sum_{c=1}^C TP_c}{\sum_{c=1}^C \text{Support}_c}$$

Evaluation metrics for multi-class classification

Medical diagnosis in $\mathcal{C} = 3$ classes: {Healthy, Cold, Flue}

100 test results:

True label	Predicted label			Support of \mathbf{c}
	c	Healthy	Cold	
	Healthy	60	8	
	Cold	4	12	
	Flue	0	2	
				100

Precision	Recall	F_1 -Score
0.94	0.86	0.9
0.54	0.6	0.57
0.57	0.8	0.67

Unweighted: Macro average

0.68	0.75	0.71
------	------	------

Weighted average

--	--	--

Accuracy

--	--	--